

# Formula 1

Business Problem: What factors/features are needed to predict race outcome

Things to do

- Outliers
- Natural Log

Powerpoint

- Agenda
- Background
- Business Problem: What factors/features are needed to predict race outcome
- Technical Appendix

1. Import required libraries

Get Data from the erghast F1 API Contains data from all the races since 1950

```
In [ ]: import pandas as pd
import os
import matplotlib.pyplot as plt
import numpy as np
import statsmodels.api as sm
from sklearn.linear_model import LinearRegression

from statsmodels.formula.api import ols
from scipy.stats import norm, probplot
from sklearn.preprocessing import StandardScaler

db_dir = os.getcwd()
pit_stops = pd.read_csv(db_dir + r'/data/pit_stops.csv')
races = pd.read_csv(db_dir + r'/data/races.csv')
drivers = pd.read_csv(db_dir + r'/data/drivers.csv')
results = pd.read_csv(db_dir + r'/data/results.csv')
driver_standings = pd.read_csv(db_dir + r'/data/driver_standings.csv')
circuits = pd.read_csv(db_dir + r'/data/circuits.csv')
qualifying = pd.read_csv(db_dir + r'/data/qualifying.csv')

merged_df = pd.merge(races, pit_stops, on='raceId', how='inner')
df = pd.merge(merged_df, drivers, on='driverId', how='inner')
df = df.drop(columns=['url_x', 'url_y', 'number'])
df_new_merged = pd.merge(df, results, on=['driverId', 'raceId'], how='inner')
df_new_merged1 = pd.merge(df_new_merged, driver_standings, on=['driverId', 'raceId'], how='inner')
df_new_merged = pd.merge(df_new_merged, circuits, on=['circuitId'], how='inner')
condition1 = df_new_merged['year'] >= 2000
condition2 = df_new_merged['year'] <= 2023
df_new_merged1 = df_new_merged[condition1 & condition2 ]

df_new_merged1.to_csv('all.csv')
def convert_new_time_to_seconds_v4(time_str):
    if time_str == '\\N':
        return np.nan
    try:
        # Split by colon and reverse the list to start with seconds and go backwards to minutes, hours, etc.
        time_parts = list(map(float, time_str.split(':')[::-1]))
        total_seconds = sum(time_part * (60 ** index) for index, time_part in enumerate(time_parts))
        return total_seconds
    except ValueError:
        return np.nan

# Apply the function to the 'new_time' column
df_new_merged1['new_time_seconds'] = df_new_merged1['time'].replace('\\+', '', regex=True).apply(convert_new_time_to_seconds_v4)
df_cleaned = df_new_merged1[['raceId', 'year', 'round', 'circuitId', 'circuit_type', 'circuit_type_num', 'driverId', 'grid_pos']]
df_cleaned = df_cleaned.drop_duplicates()
df_cleaned.sort_values(by=['raceId', 'positionOrder'], inplace=True)
df_cleaned['cumulative_time_seconds'] = df_cleaned.groupby('raceId')['new_time_seconds'].cumsum()

# Creating a variable for each of the circuit types
df_cleaned['Race_Circuit'] = np.where(df_cleaned['circuit_type_num']== 3, 1, 0)
df_cleaned['Street_Circuit'] = np.where(df_cleaned['circuit_type_num']== 2, 1, 0)
```

```
df_cleaned['Road_Circuit'] = np.where(df_cleaned['circuit_type_num']== 1, 1, 0)
df_cleaned['rank'] = df_cleaned['rank'].astype(int)

#Dropping NULL values in Cumulative race finish times - Racers not able to finish race
df_cleaned = df_cleaned.fillna(0)
print(df_new_merged1.columns)
print(df_cleaned.columns)
df_cleaned.to_csv('clean.csv')
```

```
Index(['raceId', 'year', 'round', 'circuitId', 'name_x', 'date', 'time_x',
      'fp1_date', 'fp1_time', 'fp2_date', 'fp2_time', 'fp3_date', 'fp3_time',
      'quali_date', 'quali_time', 'sprint_date', 'sprint_time', 'driverId',
      'stop', 'lap', 'time_y', 'duration', 'milliseconds_x', 'driverRef',
      'code', 'forename', 'surname', 'dob', 'nationality', 'resultId',
      'constructorId', 'number', 'grid', 'position', 'positionText',
      'positionOrder', 'points', 'laps', 'time', 'milliseconds_y',
      'fastestLap', 'rank', 'fastestLapTime', 'fastestLapSpeed', 'statusId',
      'circuitRef', 'name_y', 'location', 'country', 'lat', 'lng', 'alt',
      'url', 'circuit_type', 'circuit_type_num', 'new_time_seconds'],
      dtype='object')
Index(['raceId', 'year', 'round', 'circuitId', 'circuit_type',
      'circuit_type_num', 'driverId', 'grid', 'rank', 'position',
      'positionOrder', 'time', 'new_time_seconds', 'stop', 'laps',
      'constructorId', 'points', 'cumulative_time_seconds', 'Race_Circuit',
      'Street_Circuit', 'Road_Circuit'],
      dtype='object')
```

/var/folders/mr/l3x9p9wd385gcyns864s9f2w0000gn/T/ipykernel\_3325/1563000677.py:45: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_new_merged1['new_time_seconds'] = df_new_merged1['time'].replace('\+', '', regex=True).apply(convert_new_time_t  
o_seconds_v4)
```

```
In [ ]: import pandas as pd
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from statsmodels.stats.diagnostic import het_breuschpagan
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Assuming df_cleaned is your DataFrame
Xinitial = df_cleaned[['circuitId', 'driverId', 'grid', 'stop', 'rank', 'Race_Circuit', 'Street_Circuit', 'laps', 'y  
yinitial = df_cleaned['cumulative_time_seconds']]

# Splitting the data
X_train, X_test, y_train, y_test = train_test_split(Xinitial, yinitial, train_size=0.75, random_state=0)
X_train = sm.add_constant(X_train)
X_test = sm.add_constant(X_test)

# Fit the OLS model
model = sm.OLS(y_train, X_train)
results = model.fit()
print("OLS Model Summary:")
print(results.summary())

# To Handle multicollinearity
vif_data = pd.DataFrame()
vif_data["feature"] = X_train.columns
vif_data["VIF"] = [variance_inflation_factor(X_train.values, i) for i in range(X_train.shape[1])]
print(vif_data)

# Predicting on the training data
predicted_y = results.predict(X_train)
residuals = y_train - predicted_y

# Plotting the residuals
fig, ax = plt.subplots(1, 3, figsize=(18, 5))
# Histogram of residuals
sns.histplot(residuals, kde=True, ax=ax[0])
ax[0].set_title('Histogram of Residuals')
ax[0].set_xlabel('Residuals')
ax[0].set_ylabel('Frequency')
# Q-Q plot of residuals
stats.probplot(residuals, dist="norm", plot=ax[1])
ax[1].set_title('Q-Q Plot of Residuals')
# Residuals vs. fitted values
```

```

ax[2].scatter(results.fittedvalues, residuals)
ax[2].axhline(0, color='red', linestyle='dashed', linewidth=2)
ax[2].set_title('Residuals vs. Fitted Values')
ax[2].set_xlabel('Fitted Values')
ax[2].set_ylabel('Residuals')
plt.tight_layout()
plt.show()

# Perform the Breusch-Pagan Test
bp_test = het_breuschpagan(results.resid, results.model.exog)
measures = ('LM Statistic', 'LM-Test p-value', 'F-Statistic', 'F-Test p-value')
bp_results = dict(zip(measures, bp_test))
print("Breusch-Pagan Test Results:")
print(bp_results)
if bp_results['LM-Test p-value'] < 0.05 or bp_results['F-Test p-value'] < 0.05:
    print('The data has heteroskedasticity (p-value < 0.05). Applying Weighted Least Squares (WLS).')
    # Treat heteroskedasticity with WLS
    weights = 1 / (results.resid ** 2)
    wls_model = sm.WLS(y_train, X_train, weights=weights)
    wls_results = wls_model.fit()
    print("WLS Model Summary:")
    print(wls_results.summary())
else:
    print('The data does not have significant heteroskedasticity (p-value >= 0.05). No need to apply WLS.')

# Fit the OLS model with robust standard errors (regardless of heteroskedasticity result for comparison)
results_robust = results.get_robustcov_results(cov_type='HC3')
print("OLS Model with Robust Standard Errors Summary:")
print(results_robust.summary())

```

OLS Model Summary:


OLS Regression Results						
=====						
Dep. Variable:	cumulative_time_seconds	R-squared:	0.220			
Model:	OLS	Adj. R-squared:	0.219			
Method:	Least Squares	F-statistic:	202.2			
Date:	Thu, 01 Aug 2024	Prob (F-statistic):	0.00			
Time:	19:32:51	Log-Likelihood:	-83644.			
No. Observations:	7909	AIC:	1.673e+05			
Df Residuals:	7897	BIC:	1.674e+05			
Df Model:	11					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	-4.2e+05	6.51e+04	-6.455	0.000	-5.48e+05	-2.92e+05
circuitId	-4.9844	4.890	-1.019	0.308	-14.570	4.601
driverId	-1.0574	0.323	-3.270	0.001	-1.691	-0.424
grid	-106.8572	23.664	-4.516	0.000	-153.244	-60.470
stop	2808.3748	113.637	24.713	0.000	2585.615	3031.134
rank	-220.0211	25.548	-8.612	0.000	-270.102	-169.941
Race_Circuit	4204.1824	481.257	8.736	0.000	3260.792	5147.572
Street_Circuit	3013.6038	493.085	6.112	0.000	2047.026	3980.182
laps	112.4768	10.010	11.236	0.000	92.854	132.100
year	206.4765	32.305	6.392	0.000	143.151	269.802
constructorId	-3.5041	1.342	-2.611	0.009	-6.135	-0.874
points	263.2711	21.316	12.351	0.000	221.486	305.056
=====						
Omnibus:	4383.835	Durbin-Watson:	2.008			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	59525.424			
Skew:	2.372	Prob(JB):	0.00			
Kurtosis:	15.575	Cond. No.	1.28e+06			
=====						

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
 [2] The condition number is large, 1.28e+06. This might indicate that there are strong multicollinearity or other numerical problems.

		VIF
0	const	372204.923657
1	circuitId	1.151255
2	driverId	1.382766
3	grid	1.944453
4	stop	1.021608
5	rank	1.991964
6	Race_Circuit	4.659900
7	Street_Circuit	4.465109
8	laps	1.108966
9	year	1.419734
10	constructorId	1.146245
11	points	2.113782

```
/opt/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option_context('mode.use_inf_as_na', True):
```



## Breusch-Pagan Test Results:

{'LM Statistic': 273.1788196782718, 'LM-Test p-value': 3.845312389470348e-52, 'F-Statistic': 25.683885656759657, 'F-Test p-value': 4.224267665617925e-53}

The data has heteroskedasticity (p-value < 0.05). Applying Weighted Least Squares (WLS).

## WLS Model Summary:

## WLS Regression Results

```
=====
Dep. Variable:    cumulative_time_seconds    R-squared:                0.999
Model:                WLS    Adj. R-squared:            0.999
Method:            Least Squares    F-statistic:            6.676e+05
Date:                Thu, 01 Aug 2024    Prob (F-statistic):      0.00
Time:                19:32:52    Log-Likelihood:         -77721.
No. Observations:    7909    AIC:                    1.555e+05
Df Residuals:        7897    BIC:                    1.555e+05
Df Model:            11
Covariance Type:    nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-4.19e+05	3455.766	-121.248	0.000	-4.26e+05	-4.12e+05
circuitId	-4.7355	0.183	-25.853	0.000	-5.095	-4.376
driverId	-1.0509	0.014	-73.492	0.000	-1.079	-1.023
grid	-107.0254	1.225	-87.361	0.000	-109.427	-104.624
stop	2810.0779	6.841	410.797	0.000	2796.669	2823.487
rank	-220.5468	1.132	-194.787	0.000	-222.766	-218.327
Race_Circuit	4274.1317	30.006	142.441	0.000	4215.311	4332.952
Street_Circuit	3078.4677	31.167	98.773	0.000	3017.372	3139.563
laps	112.8767	0.394	286.160	0.000	112.104	113.650
year	205.9096	1.710	120.450	0.000	202.559	209.261
constructorId	-3.4535	0.050	-69.276	0.000	-3.551	-3.356
points	265.3157	1.066	248.853	0.000	263.226	267.406

```
=====
Omnibus:                31156.029    Durbin-Watson:            1.905
Prob(Omnibus):            0.000    Jarque-Bera (JB):        1321.173
Skew:                    0.424    Prob(JB):                1.29e-287
Kurtosis:                1.186    Cond. No.:                4.10e+06
=====
```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.1e+06. This might indicate that there are strong multicollinearity or other numerical problems.

## OLS Model with Robust Standard Errors Summary:

## OLS Regression Results

```
=====
Dep. Variable:    cumulative_time_seconds    R-squared:                0.220
Model:                OLS    Adj. R-squared:            0.219
Method:            Least Squares    F-statistic:            220.4
Date:                Thu, 01 Aug 2024    Prob (F-statistic):      0.00
Time:                19:32:52    Log-Likelihood:         -83644.
No. Observations:    7909    AIC:                    1.673e+05
Df Residuals:        7897    BIC:                    1.674e+05
Df Model:            11
Covariance Type:    HC3
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-4.2e+05	6.89e+04	-6.099	0.000	-5.55e+05	-2.85e+05
circuitId	-4.9844	4.541	-1.098	0.272	-13.885	3.916
driverId	-1.0574	0.309	-3.423	0.001	-1.663	-0.452
grid	-106.8572	26.367	-4.053	0.000	-158.543	-55.172
stop	2808.3748	157.227	17.862	0.000	2500.168	3116.582
rank	-220.0211	23.649	-9.303	0.000	-266.380	-173.662
Race_Circuit	4204.1824	378.752	11.100	0.000	3461.728	4946.637
Street_Circuit	3013.6038	360.517	8.359	0.000	2306.895	3720.313
laps	112.4768	9.931	11.325	0.000	93.009	131.945
year	206.4765	34.126	6.050	0.000	139.580	273.373
constructorId	-3.5041	1.306	-2.682	0.007	-6.065	-0.943
points	263.2711	21.935	12.002	0.000	220.272	306.270

```
=====
Omnibus:                4383.835    Durbin-Watson:            2.008
Prob(Omnibus):            0.000    Jarque-Bera (JB):        59525.424
Skew:                    2.372    Prob(JB):                0.00
Kurtosis:                15.575    Cond. No.:                1.28e+06
=====
```

## Notes:

[1] Standard Errors are heteroscedasticity robust (HC3)

[2] The condition number is large, 1.28e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [ ]: # Hypothesis Testing for the different circuit: Road, Race and Street

hypothesis = '(Race_Circuit=0, Street_Circuit=0)'
#Pass the hypothesis to the Wald_Test
print(results_robust.wald_test(hypothesis))

#We do not embrace the null and there is a difference between the different types of track and the time it takes to

<F test: F=array([[62.26236896]]), p=1.4817501629469654e-27, df_denom=7.9e+03, df_num=2>
/opt/anaconda3/lib/python3.11/site-packages/statsmodels/base/model.py:1906: FutureWarning: The behavior of wald_test
will change after 0.14 to returning scalar test statistic values. To get the future behavior now, set scalar to True.
To silence this message while retaining the legacy behavior, set scalar to False.
  warnings.warn(
```

```
In [ ]: import pandas as pd
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from statsmodels.stats.diagnostic import het_breuschpagan
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Assuming df_cleaned is your DataFrame

# Creating Dummy variables for splitting the two different types of circuits (street and race)
df_cleaned['Race_Circuit_grid'] = df_cleaned['Race_Circuit'] * df_cleaned['grid']
df_cleaned['Street_Circuit_grid'] = df_cleaned['Street_Circuit'] * df_cleaned['grid']

# Define the independent variables (X) and the dependent variable (y)
Xinitial = df_cleaned[['grid', 'stop', 'rank', 'Race_Circuit', 'Street_Circuit', 'laps', 'year', 'Race_Circuit_grid']]
yinitial = df_cleaned['cumulative_time_seconds']

# Apply Natural Log Transformation to the dependent variable
yinitial_transformed = np.log(df_cleaned['cumulative_time_seconds'] + 1)
print(f"Skewness of transformed cumulative_time_seconds: {yinitial_transformed.skew()}")

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(Xinitial, yinitial_transformed, train_size=0.75, random_state=0)

# Add a constant to the model (intercept)
X_train = sm.add_constant(X_train)
X_test = sm.add_constant(X_test)

# Calculate VIF for each predictor
vif_data = pd.DataFrame()
vif_data["feature"] = X_train.columns
vif_data["VIF"] = [variance_inflation_factor(X_train.values, i) for i in range(X_train.shape[1])]
print("VIF Data:")
print(vif_data)

# Fit the OLS model
model = sm.OLS(y_train, X_train)
results = model.fit()
print("OLS Model Summary with Transformed Dependent Variable:")
print(results.summary())

# Predicting on the training data
predicted_y = results.predict(X_train)
residuals = y_train - predicted_y

# Plotting the residuals
fig, ax = plt.subplots(1, 3, figsize=(18, 5))
# Histogram of residuals
sns.histplot(residuals, kde=True, ax=ax[0])
ax[0].set_title('Histogram of Residuals')
ax[0].set_xlabel('Residuals')
ax[0].set_ylabel('Frequency')
# Q-Q plot of residuals
stats.probplot(residuals, dist="norm", plot=ax[1])
ax[1].set_title('Q-Q Plot of Residuals')
# Residuals vs. fitted values
ax[2].scatter(results.fittedvalues, residuals)
ax[2].axhline(0, color='red', linestyle='dashed', linewidth=2)
ax[2].set_title('Residuals vs. Fitted Values')
ax[2].set_xlabel('Fitted Values')
ax[2].set_ylabel('Residuals')
```

```

plt.tight_layout()
plt.show()

# Perform the Breusch-Pagan Test
from statsmodels.stats.diagnostic import het_breuschpagan
bp_test = het_breuschpagan(results.resid, results.model.exog)
measures = ('LM Statistic', 'LM-Test p-value', 'F-Statistic', 'F-Test p-value')
bp_results = dict(zip(measures, bp_test))
print("Breusch-Pagan Test Results:")
print(bp_results)

# Apply Weighted Least Squares (WLS) if heteroskedasticity is detected
if bp_results['LM-Test p-value'] < 0.05 or bp_results['F-Test p-value'] < 0.05:
    print('The data has heteroskedasticity (p-value < 0.05). Applying Weighted Least Squares (WLS).')
    weights = 1 / (results.resid ** 2)
    wls_model = sm.WLS(y_train, X_train, weights=weights)
    wls_results = wls_model.fit()
    print("WLS Model Summary:")
    print(wls_results.summary())
    # Predicting on the test data using the WLS model
    y_pred_train_wls = wls_results.predict(X_train)
    y_pred_test_wls = wls_results.predict(X_test)
    # Perform RMSE and R2 test with Wald test on train data
    rmse_train_wls = np.sqrt(mean_squared_error(y_train, y_pred_train_wls))
    r2_train_wls = r2_score(y_train, y_pred_train_wls)
    # Perform RMSE and R2 test with Wald test on test data
    rmse_test_wls = np.sqrt(mean_squared_error(y_test, y_pred_test_wls))
    r2_test_wls = r2_score(y_test, y_pred_test_wls)

    # Output the WLS results
    print(f"Train RMSE (WLS): {rmse_train_wls}")
    print(f"Train R-squared (WLS): {r2_train_wls}")
    print(" ")
    print(f"Test RMSE (WLS): {rmse_test_wls}")
    print(f"Test R-squared (WLS): {r2_test_wls}")
else:
    print('The data does not have significant heteroskedasticity (p-value >= 0.05). No need to apply WLS.')

```

Skewness of transformed cumulative\_time\_seconds: -0.19598903775263365

VIF Data:

	feature	VIF
0	const	270290.818699
1	grid	18.445503
2	stop	1.014493
3	rank	1.593518
4	Race_Circuit	18.940374
5	Street_Circuit	18.810190
6	laps	1.101929
7	year	1.029900
8	Race_Circuit_grid	25.621225
9	Street_Circuit_grid	20.115156

OLS Model Summary with Transformed Dependent Variable:

OLS Regression Results

```

=====
Dep. Variable:    cumulative_time_seconds    R-squared:                0.320
Model:            OLS                      Adj. R-squared:           0.320
Method:           Least Squares             F-statistic:             413.8
Date:             Fri, 02 Aug 2024           Prob (F-statistic):      0.00
Time:             11:54:55                  Log-Likelihood:         -21975.
No. Observations: 7909                     AIC:                    4.397e+04
Df Residuals:     7899                     BIC:                    4.404e+04
Df Model:         9
Covariance Type:  nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-52.7075	22.782	-2.314	0.021	-97.367	-8.048
grid	-0.3255	0.030	-10.872	0.000	-0.384	-0.267
stop	0.1843	0.047	3.961	0.000	0.093	0.275
rank	-0.2881	0.009	-30.695	0.000	-0.307	-0.270
Race_Circuit	0.0313	0.399	0.079	0.937	-0.750	0.813
Street_Circuit	0.0645	0.416	0.155	0.877	-0.751	0.880
laps	0.0330	0.004	8.059	0.000	0.025	0.041
year	0.0300	0.011	2.655	0.008	0.008	0.052
Race_Circuit_grid	0.1588	0.031	5.159	0.000	0.098	0.219
Street_Circuit_grid	0.1578	0.032	4.904	0.000	0.095	0.221

```

=====
Omnibus:         620.985    Durbin-Watson:           1.982
Prob(Omnibus):   0.000     Jarque-Bera (JB):        217.506
Skew:            -0.117    Prob(JB):                5.88e-48
Kurtosis:        2.222     Cond. No.:               1.05e+06
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.05e+06. This might indicate that there are strong multicollinearity or other numerical problems.

/opt/anaconda3/lib/python3.11/site-packages/seaborn/\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):





## Breusch-Pagan Test Results:

{'LM Statistic': 329.9577331592434, 'LM-Test p-value': 1.1355297659380292e-65, 'F-Statistic': 38.209696371501636, 'F-Test p-value': 4.062596101603407e-67}

The data has heteroskedasticity (p-value < 0.05). Applying Weighted Least Squares (WLS).

## WLS Model Summary:

## WLS Regression Results

=====						
Dep. Variable:	cumulative_time_seconds		R-squared:	1.000		
Model:	WLS		Adj. R-squared:	1.000		
Method:	Least Squares		F-statistic:	4.075e+06		
Date:	Fri, 02 Aug 2024		Prob (F-statistic):	0.00		
Time:	11:54:55		Log-Likelihood:	-17649.		
No. Observations:	7909		AIC:	3.532e+04		
Df Residuals:	7899		BIC:	3.539e+04		
Df Model:	9					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	-53.3500	0.826	-64.562	0.000	-54.970	-51.730
grid	-0.3269	0.001	-292.843	0.000	-0.329	-0.325
stop	0.1854	0.001	125.829	0.000	0.182	0.188
rank	-0.2880	0.000	-605.487	0.000	-0.289	-0.287
Race_Circuit	0.0343	0.003	11.248	0.000	0.028	0.040
Street_Circuit	0.0689	0.006	12.164	0.000	0.058	0.080
laps	0.0331	0.000	268.035	0.000	0.033	0.033
year	0.0303	0.000	74.172	0.000	0.030	0.031
Race_Circuit_grid	0.1601	0.001	151.774	0.000	0.158	0.162
Street_Circuit_grid	0.1586	0.001	138.726	0.000	0.156	0.161
=====						
Omnibus:	27486.376	Durbin-Watson:	1.975			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1315.616			
Skew:	-0.102	Prob(JB):	2.08e-286			
Kurtosis:	1.012	Cond. No.	3.86e+06			
=====						

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.86e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Train RMSE (WLS): 3.8946488009693003

Train R-squared (WLS): 0.320412471941246

Test RMSE (WLS): 3.934348436267168

Test R-squared (WLS): 0.30854996208669827